

# OS & Virtual Memory

Week 14

Suting Chen

# OS & Virtual Memory

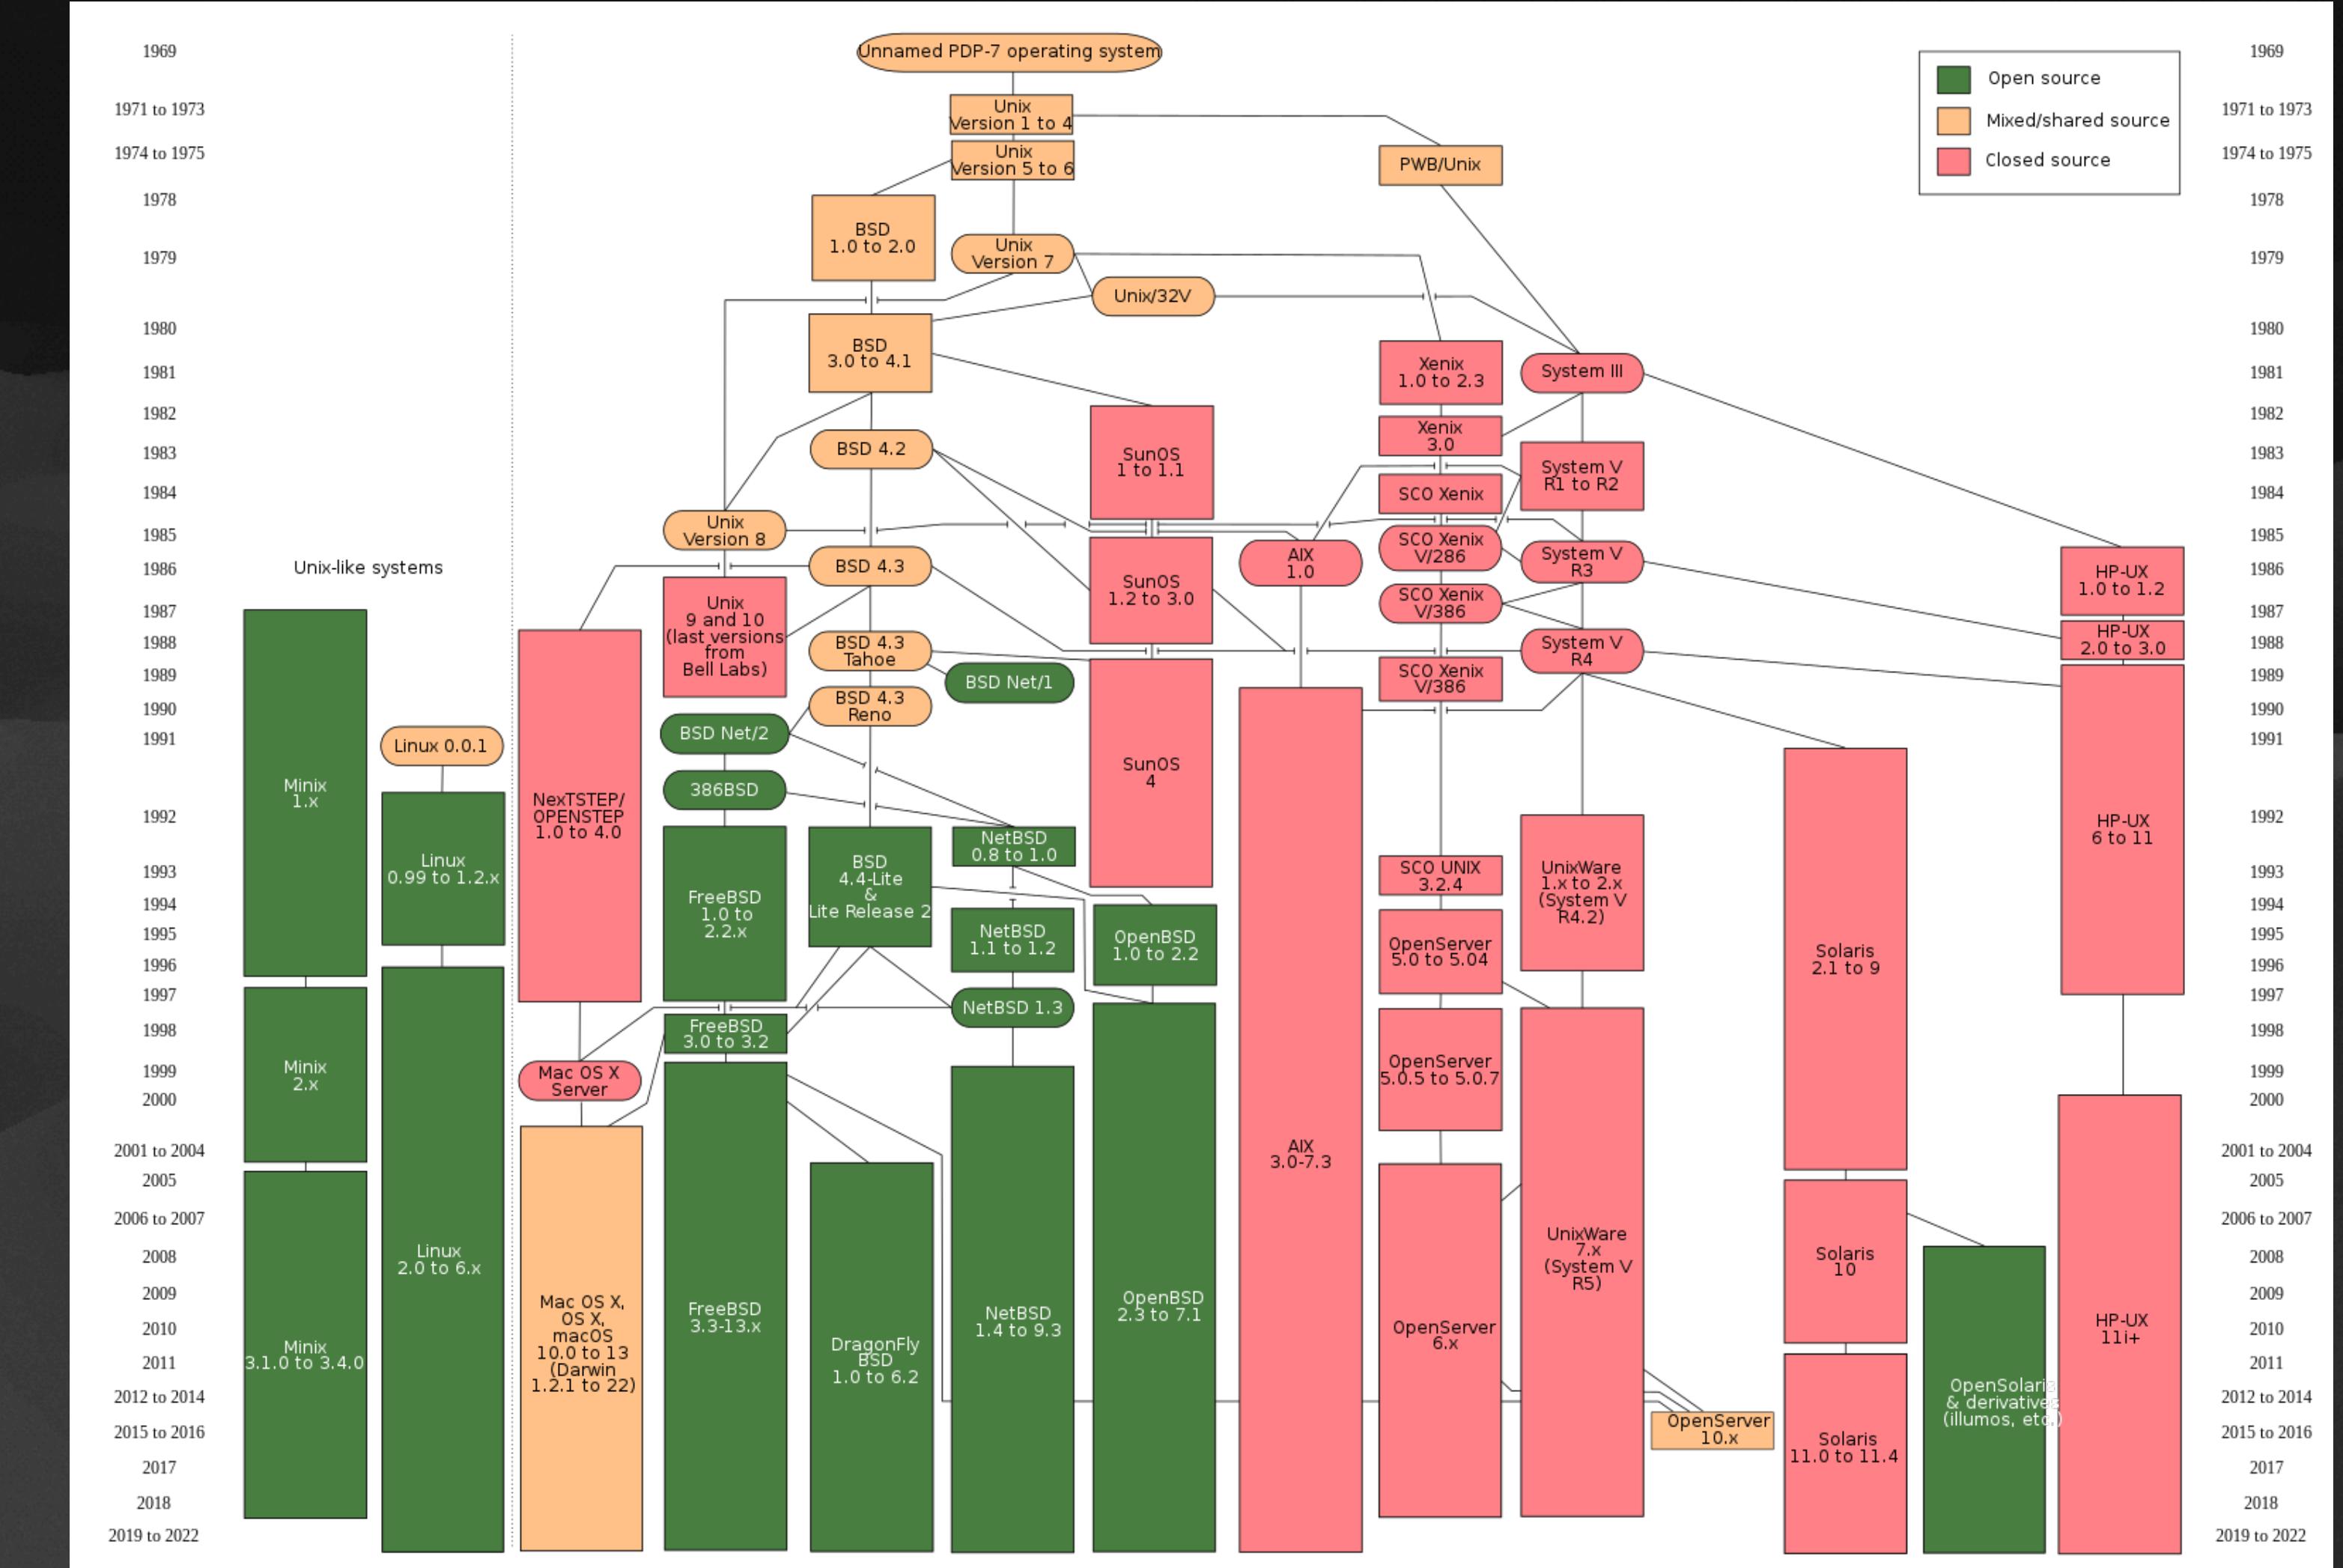
- What's Operating System?
- How OS helps your programs?
- Virtual Memory

Suting Chen

# What's Operating System?

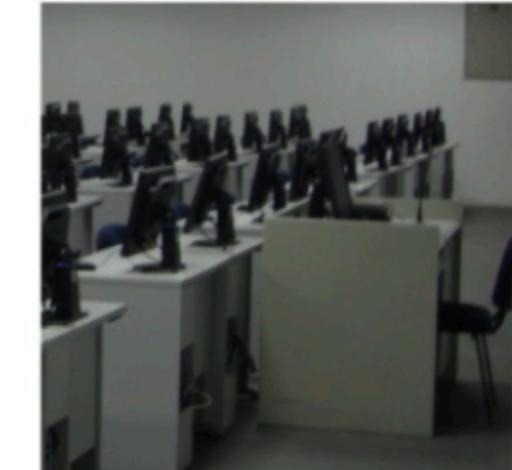
- Linux / Unix / Windows
- Debian / Fedora / Arch

Suting Chen



# What's Operating System?

Suting Chen

	目标纯粹	目标中立	目标混乱
对象纯粹	必须管理一台计算机	管理计算机硬件就行	管理什么都行
对象中立	Windows/Linux/vxWorks 当然是操作系统	固件也是操作系统	gdb也是操作系统
对象混乱	浏览器/微信/支付宝/JVM 都是操作系统	Hadoop也是操作系统	资源管理器也是操作系统
服务谁都可以			
机箱是操作系统	机房也可以算操作系统	校长为什么不是操作系统？	

# How OS helps your programs?

- Helps loading your program (CALL)
- Offers a set of functions
- Gives an illusion of the system
- Schedule your threads fairly

Suting Chen

# How OS helps your programs?

## Helps loading your program (CALL)

- Linux load a program in ELF format from disk to memory.
- Make space for stack and heap.

# How OS helps your programs?

## Offers a set of functions

- How to print to console, read from input, allocate memory?
- `printf()`, `scanf()`, `malloc()`

# How OS helps your programs?

Offers a set of functions – interacting with hardware

- printf / scanf – Keyboard
- malloc – Memory
- fopen – Disk
- socket – Network card
- pthread\_create – CPU

Suting Chen

# How OS helps your programs?

Offers a set of functions – pros

- Saves programmers' time to read docs and interact with hardware

# How OS helps your programs?

## Offers a set of functions – pros

- Saves programmers' time to read docs and interact with hardware
- Offers unified interface for multiple devices:
  - `pthread_create()` works for all platforms – Intel, AMD, Apple...
  - `fopen()` opens file on all devices – SSD, HDD and even floppy disks!

# How OS helps your programs?

## Offers a set of functions – pros

- Saves programmers' time to read docs and interact with hardware
- Offers unified interface for multiple devices:
  - `pthread_create()` works for all platforms – Intel, AMD, Apple...
  - `fopen()` opens file on all devices – SSD, HDD and even floppy disks!
- Enables permissions

```
chmod 750 /home/cs110
```

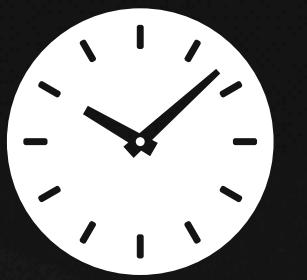
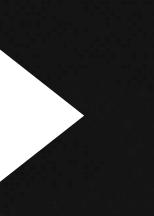
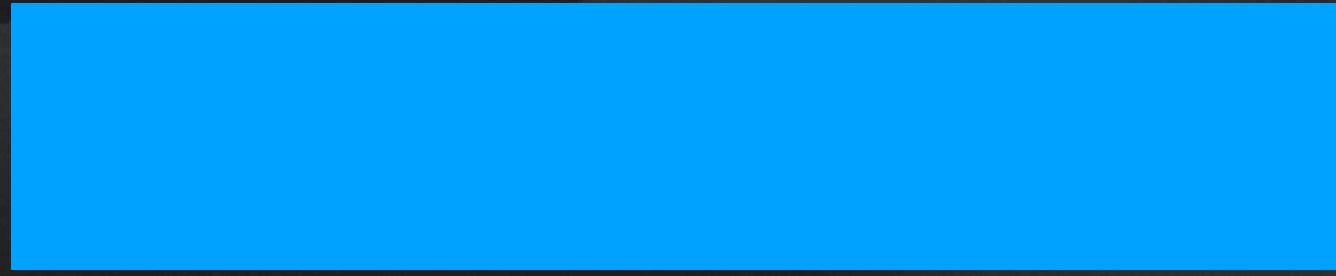
# How OS helps your programs?

**Gives an illusion of the system**

- A real CPU
- Ideal memory addresses

# Gives an illusion of the system – A real CPU

When running single program...



# Gives an illusion of the system – A real CPU

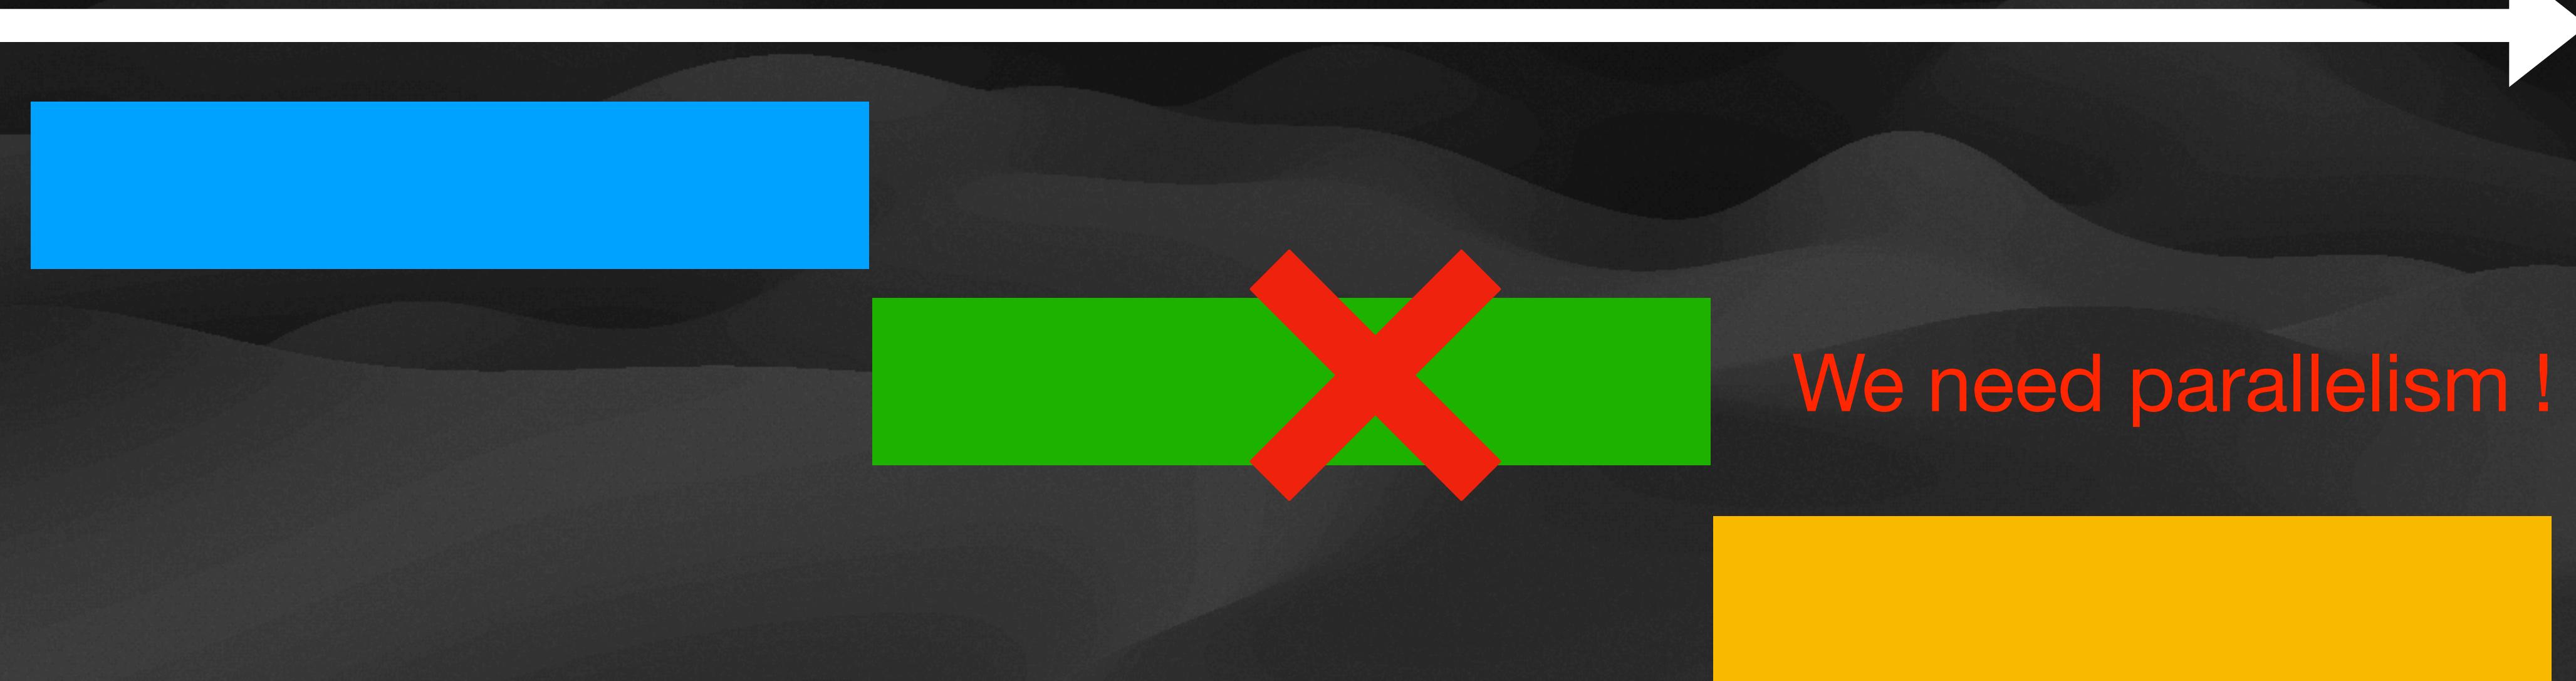
When running multiple programs...



Our naive CPU can only execute  
one instruction at a time

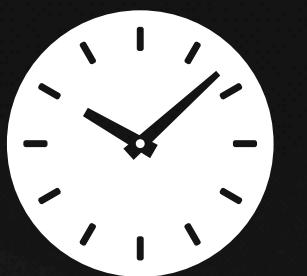
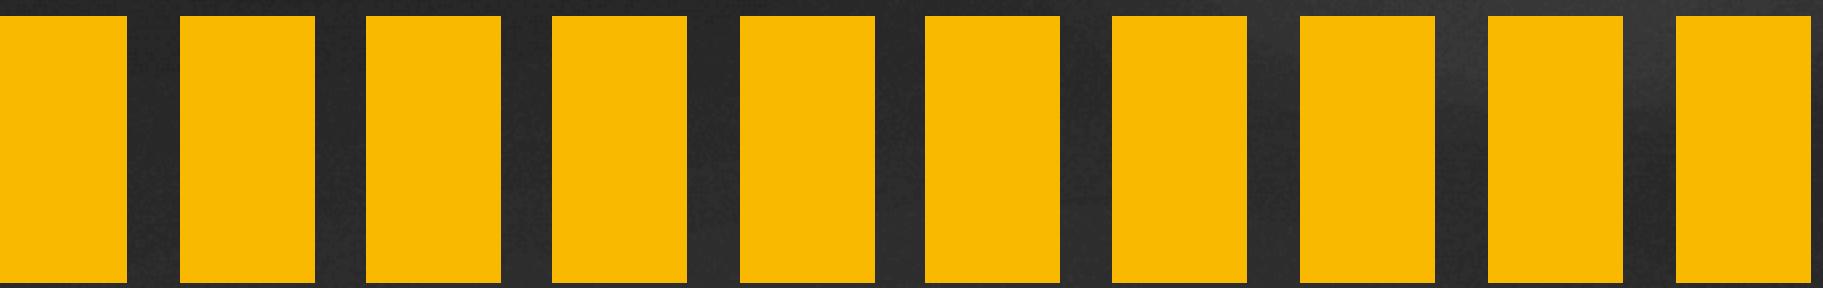
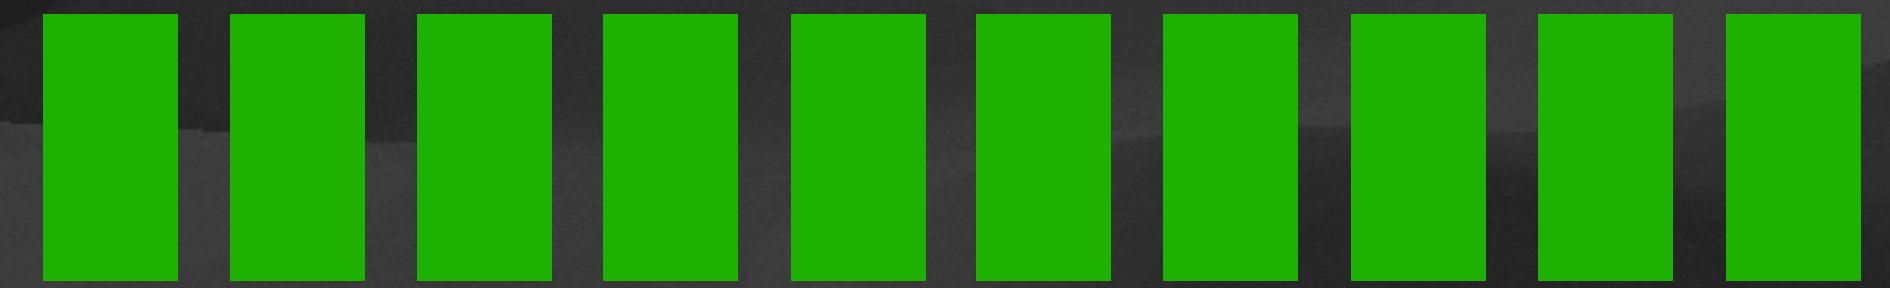
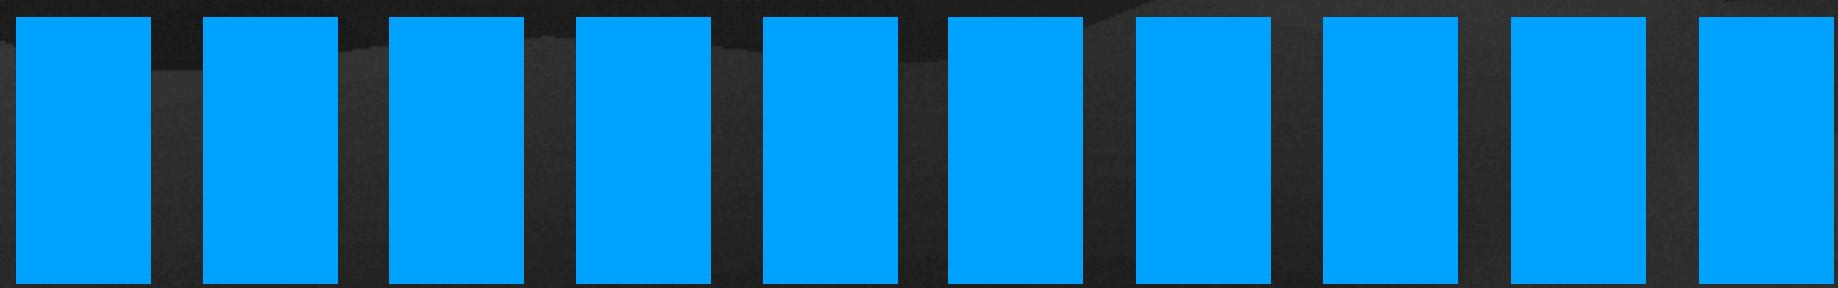
# Gives an illusion of the system – A real CPU

When running multiple programs...



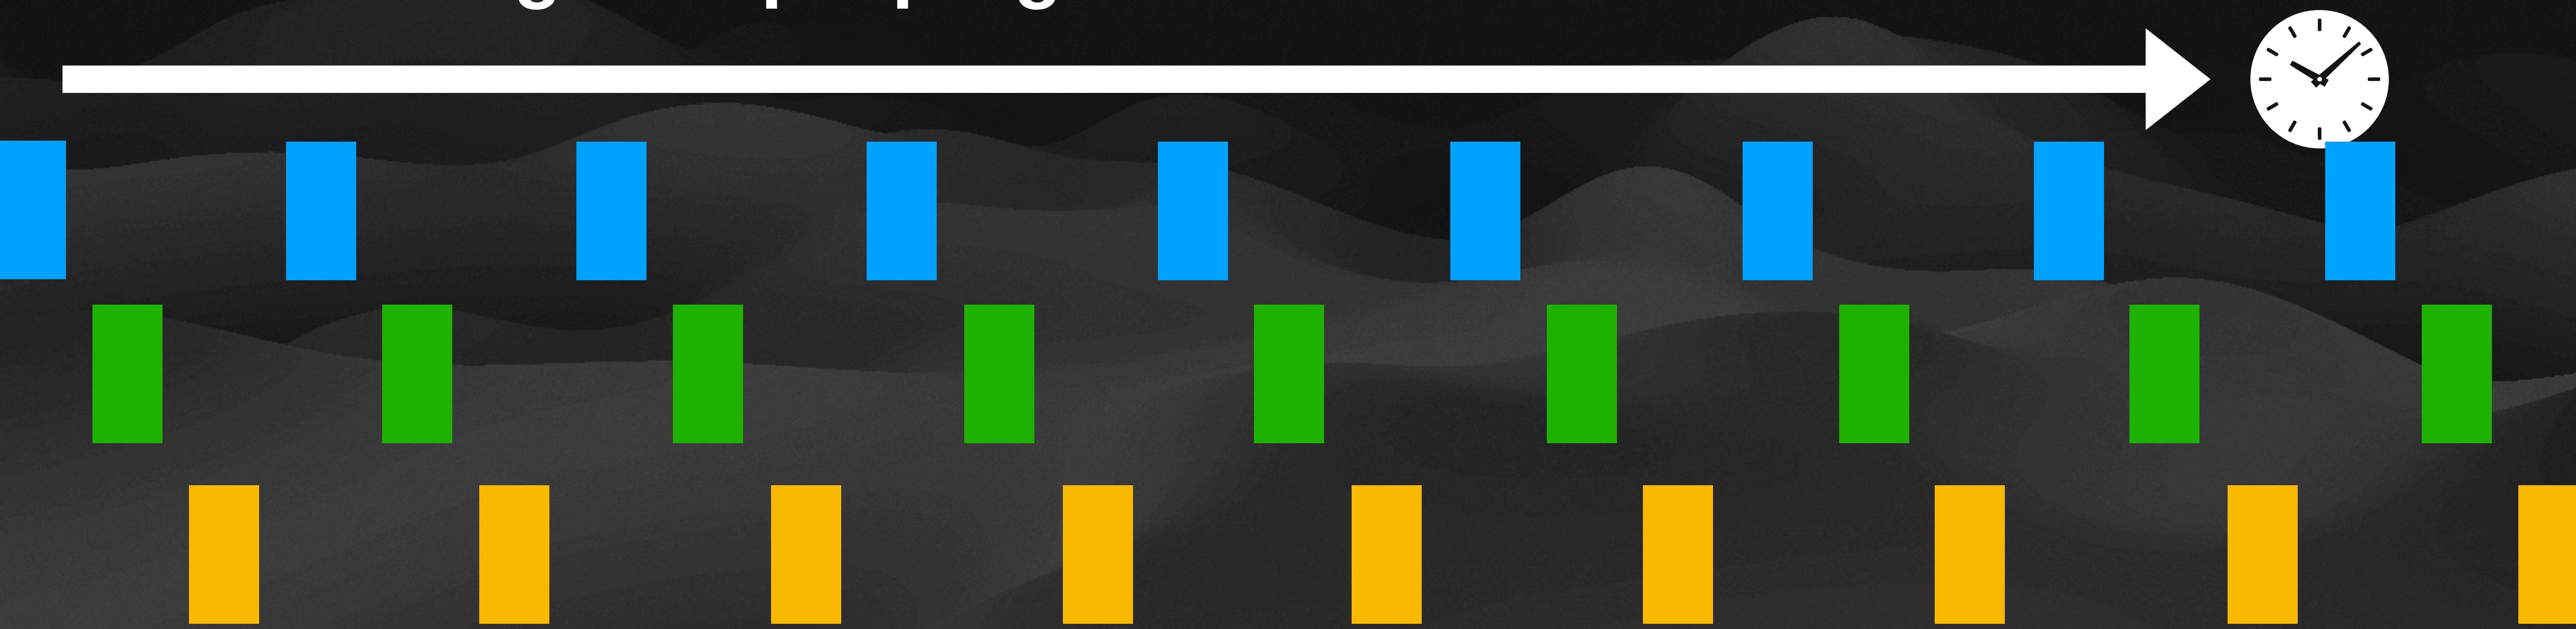
# Gives an illusion of the system – A real CPU

When running multiple programs...



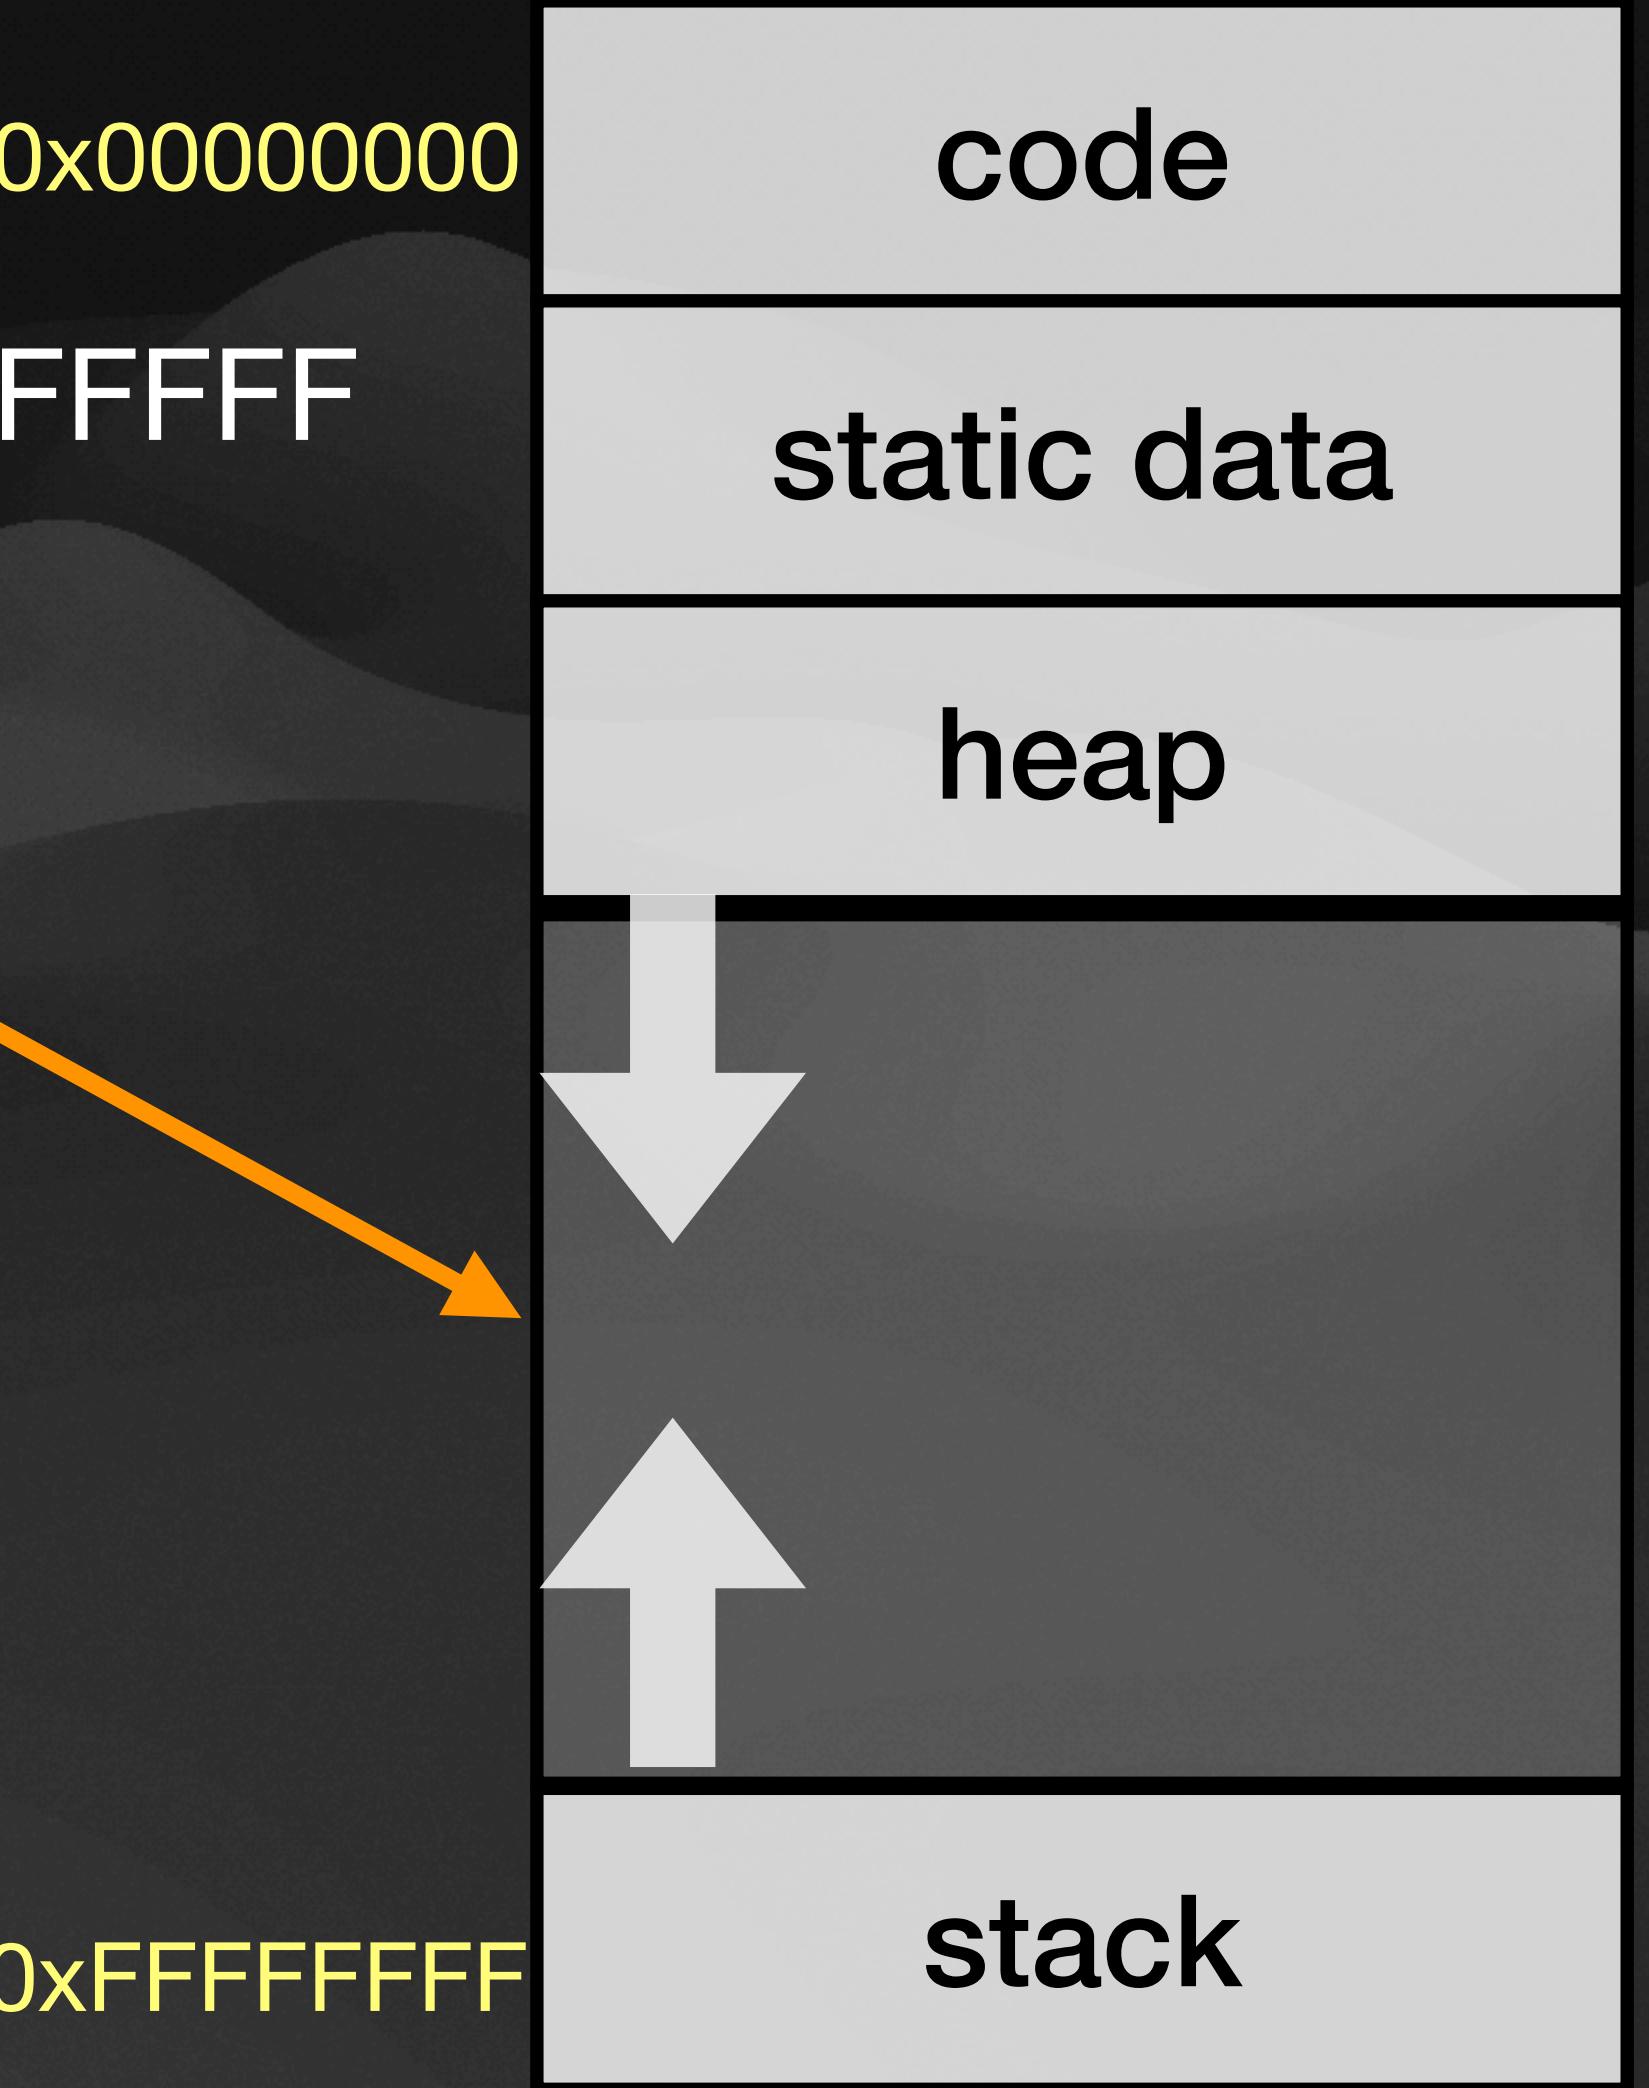
# Gives an illusion of the system – A real CPU

When running multiple programs...



# Gives an illusion of the system – Ideal memory addresses

- Every program assert it has memory size of 0xFFFFFFFF
- Problem:
  - We don't have so much memory to **waste**
  - We only have one 0x0 and one 0xFFFFFFFF



# Gives an illusion of the system – Ideal memory addresses

If you have a 64-bit machine (certainly you do!)

- 64-bit addresses!
- Total size  $2^{64}$  Bytes = 16 EB =  $16 \times 1024$  PB
  - Impossible to find such memory (at least in 2023)

Suting Chen

实际上市面上的64位机器一般不会真的有64位的内存空间，这里为了演示方便忽略了很多问题  
具体可以去看AMD64 Architecture Programmer's Manual, 你也可以用`lscpu`查看自己的CPU情况  
大部分情况下都是48位的Virtual Address Space (这里当故事听就好)

# Virtual Memory

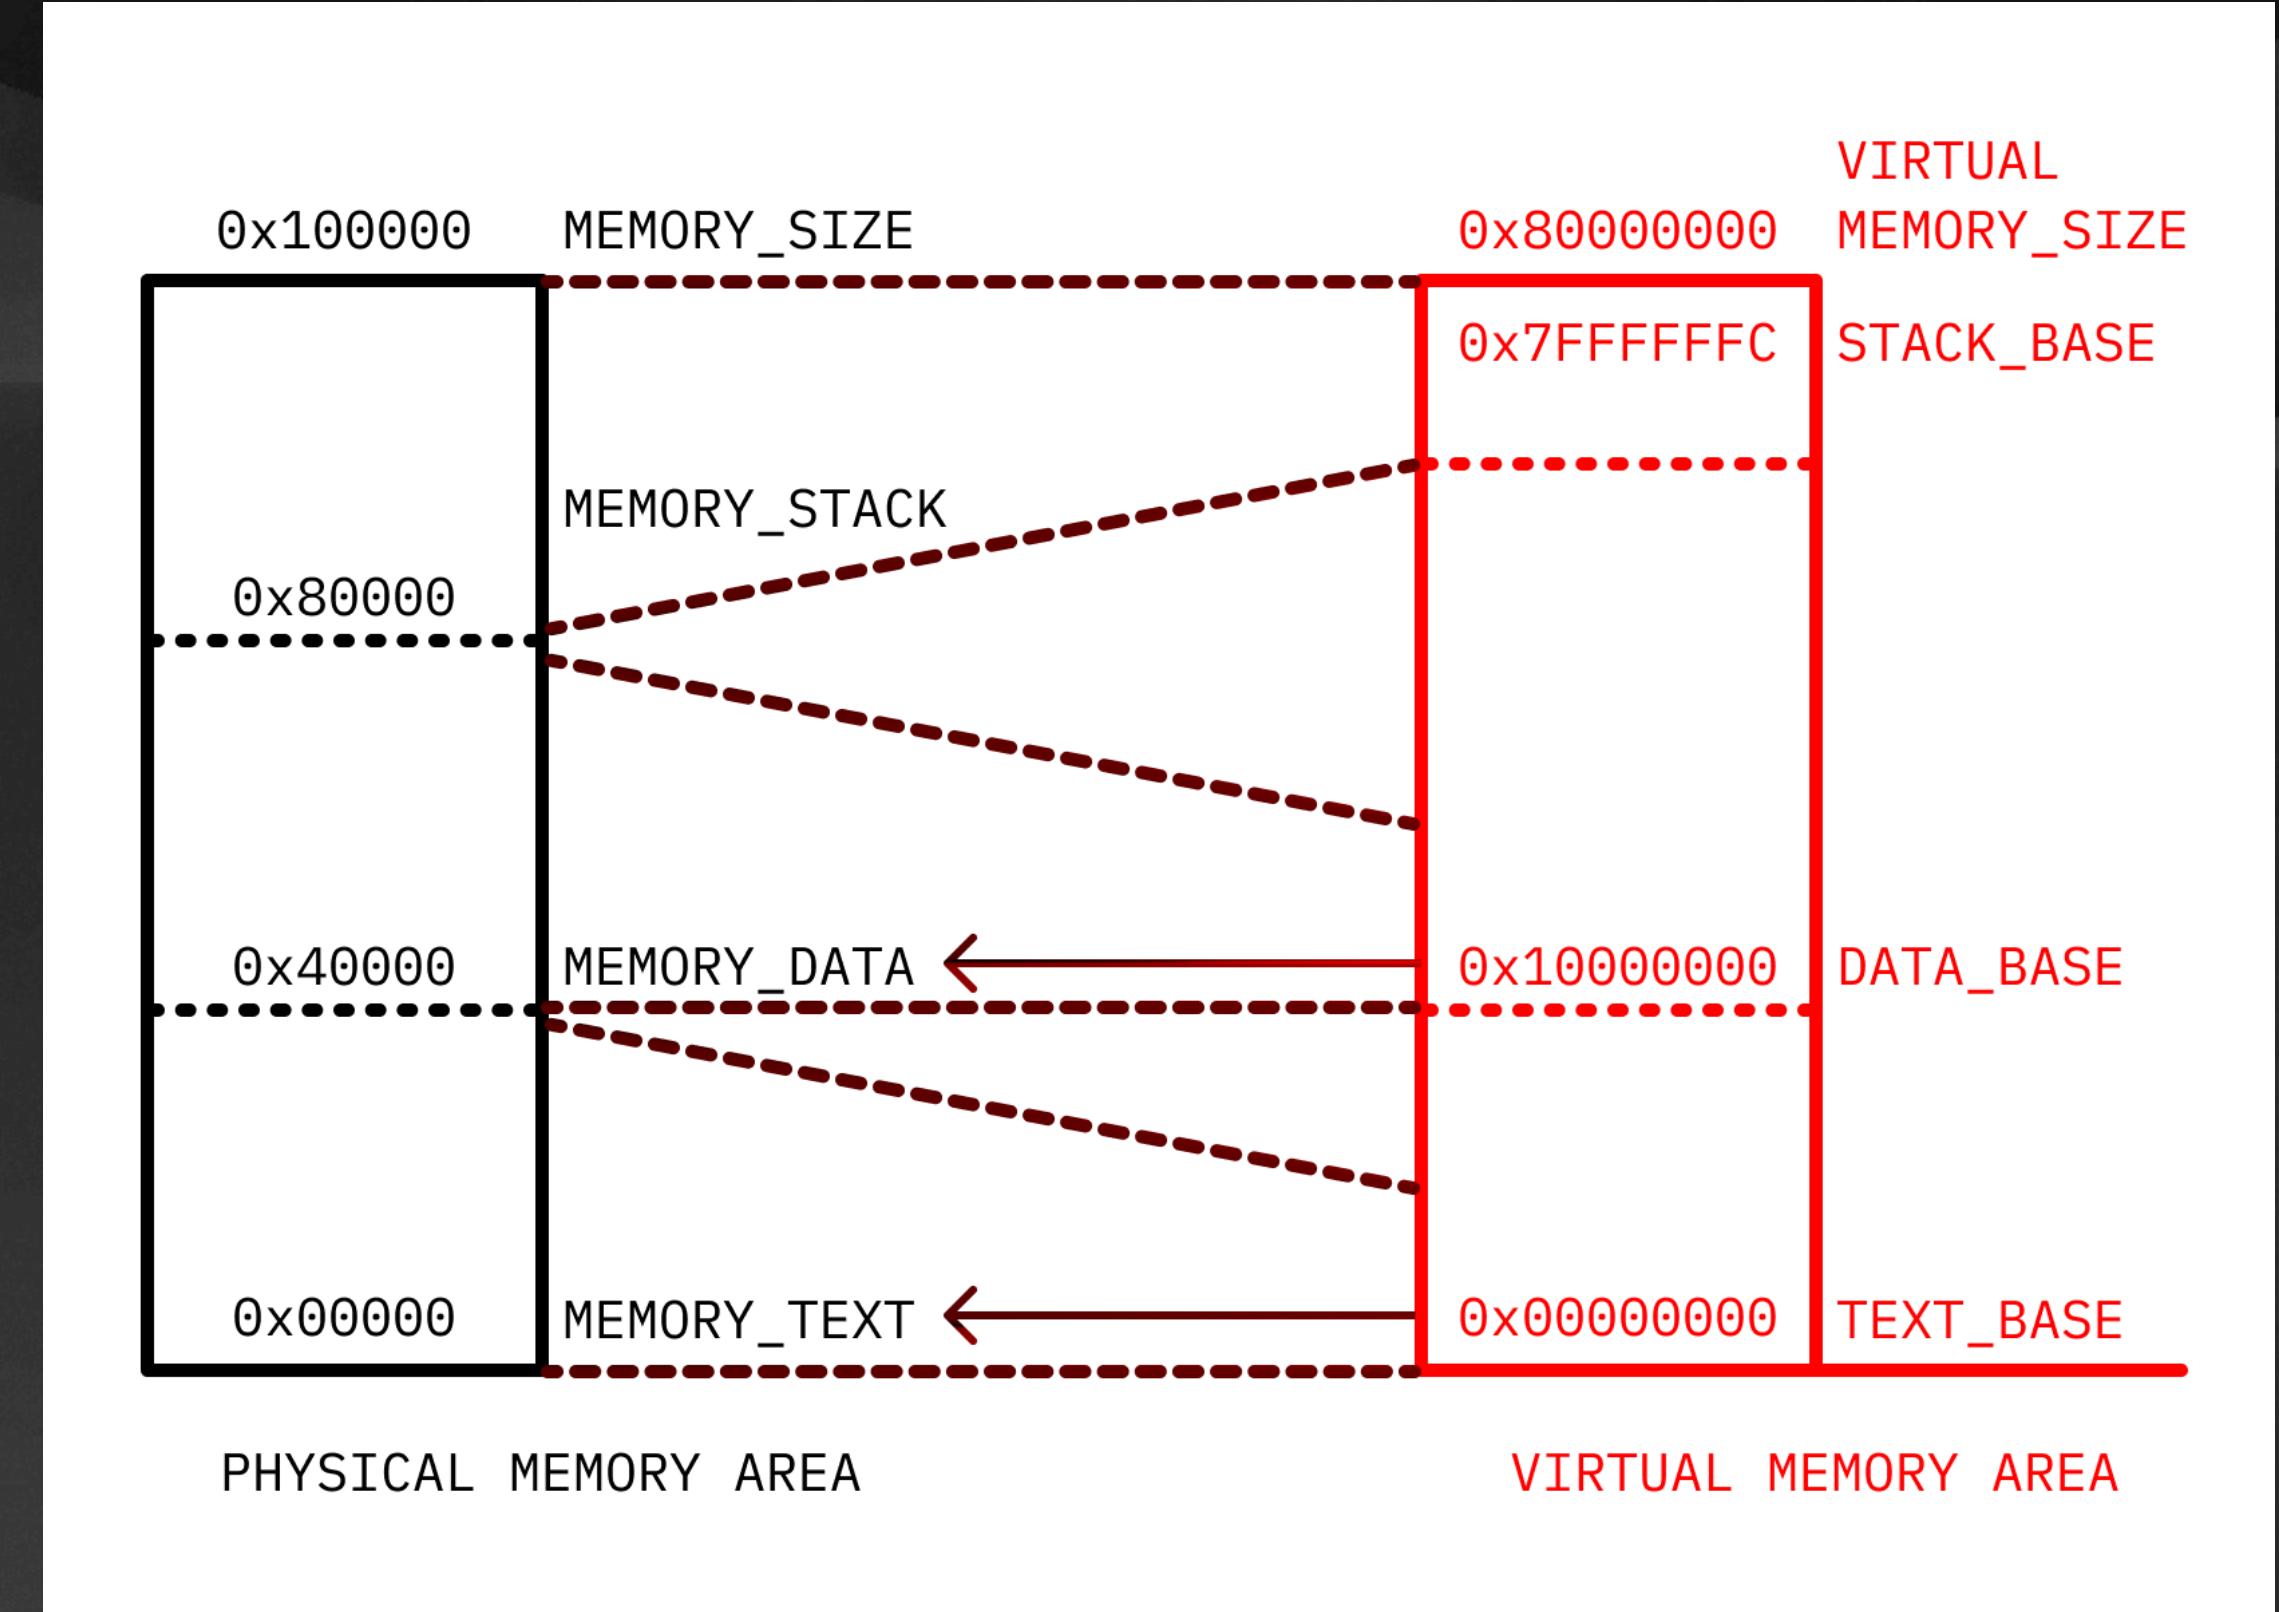
## A very modern solution

- Idea: hide physical address from programs

- Introduce **Virtual Address Space**

- Project 1.2 ➔

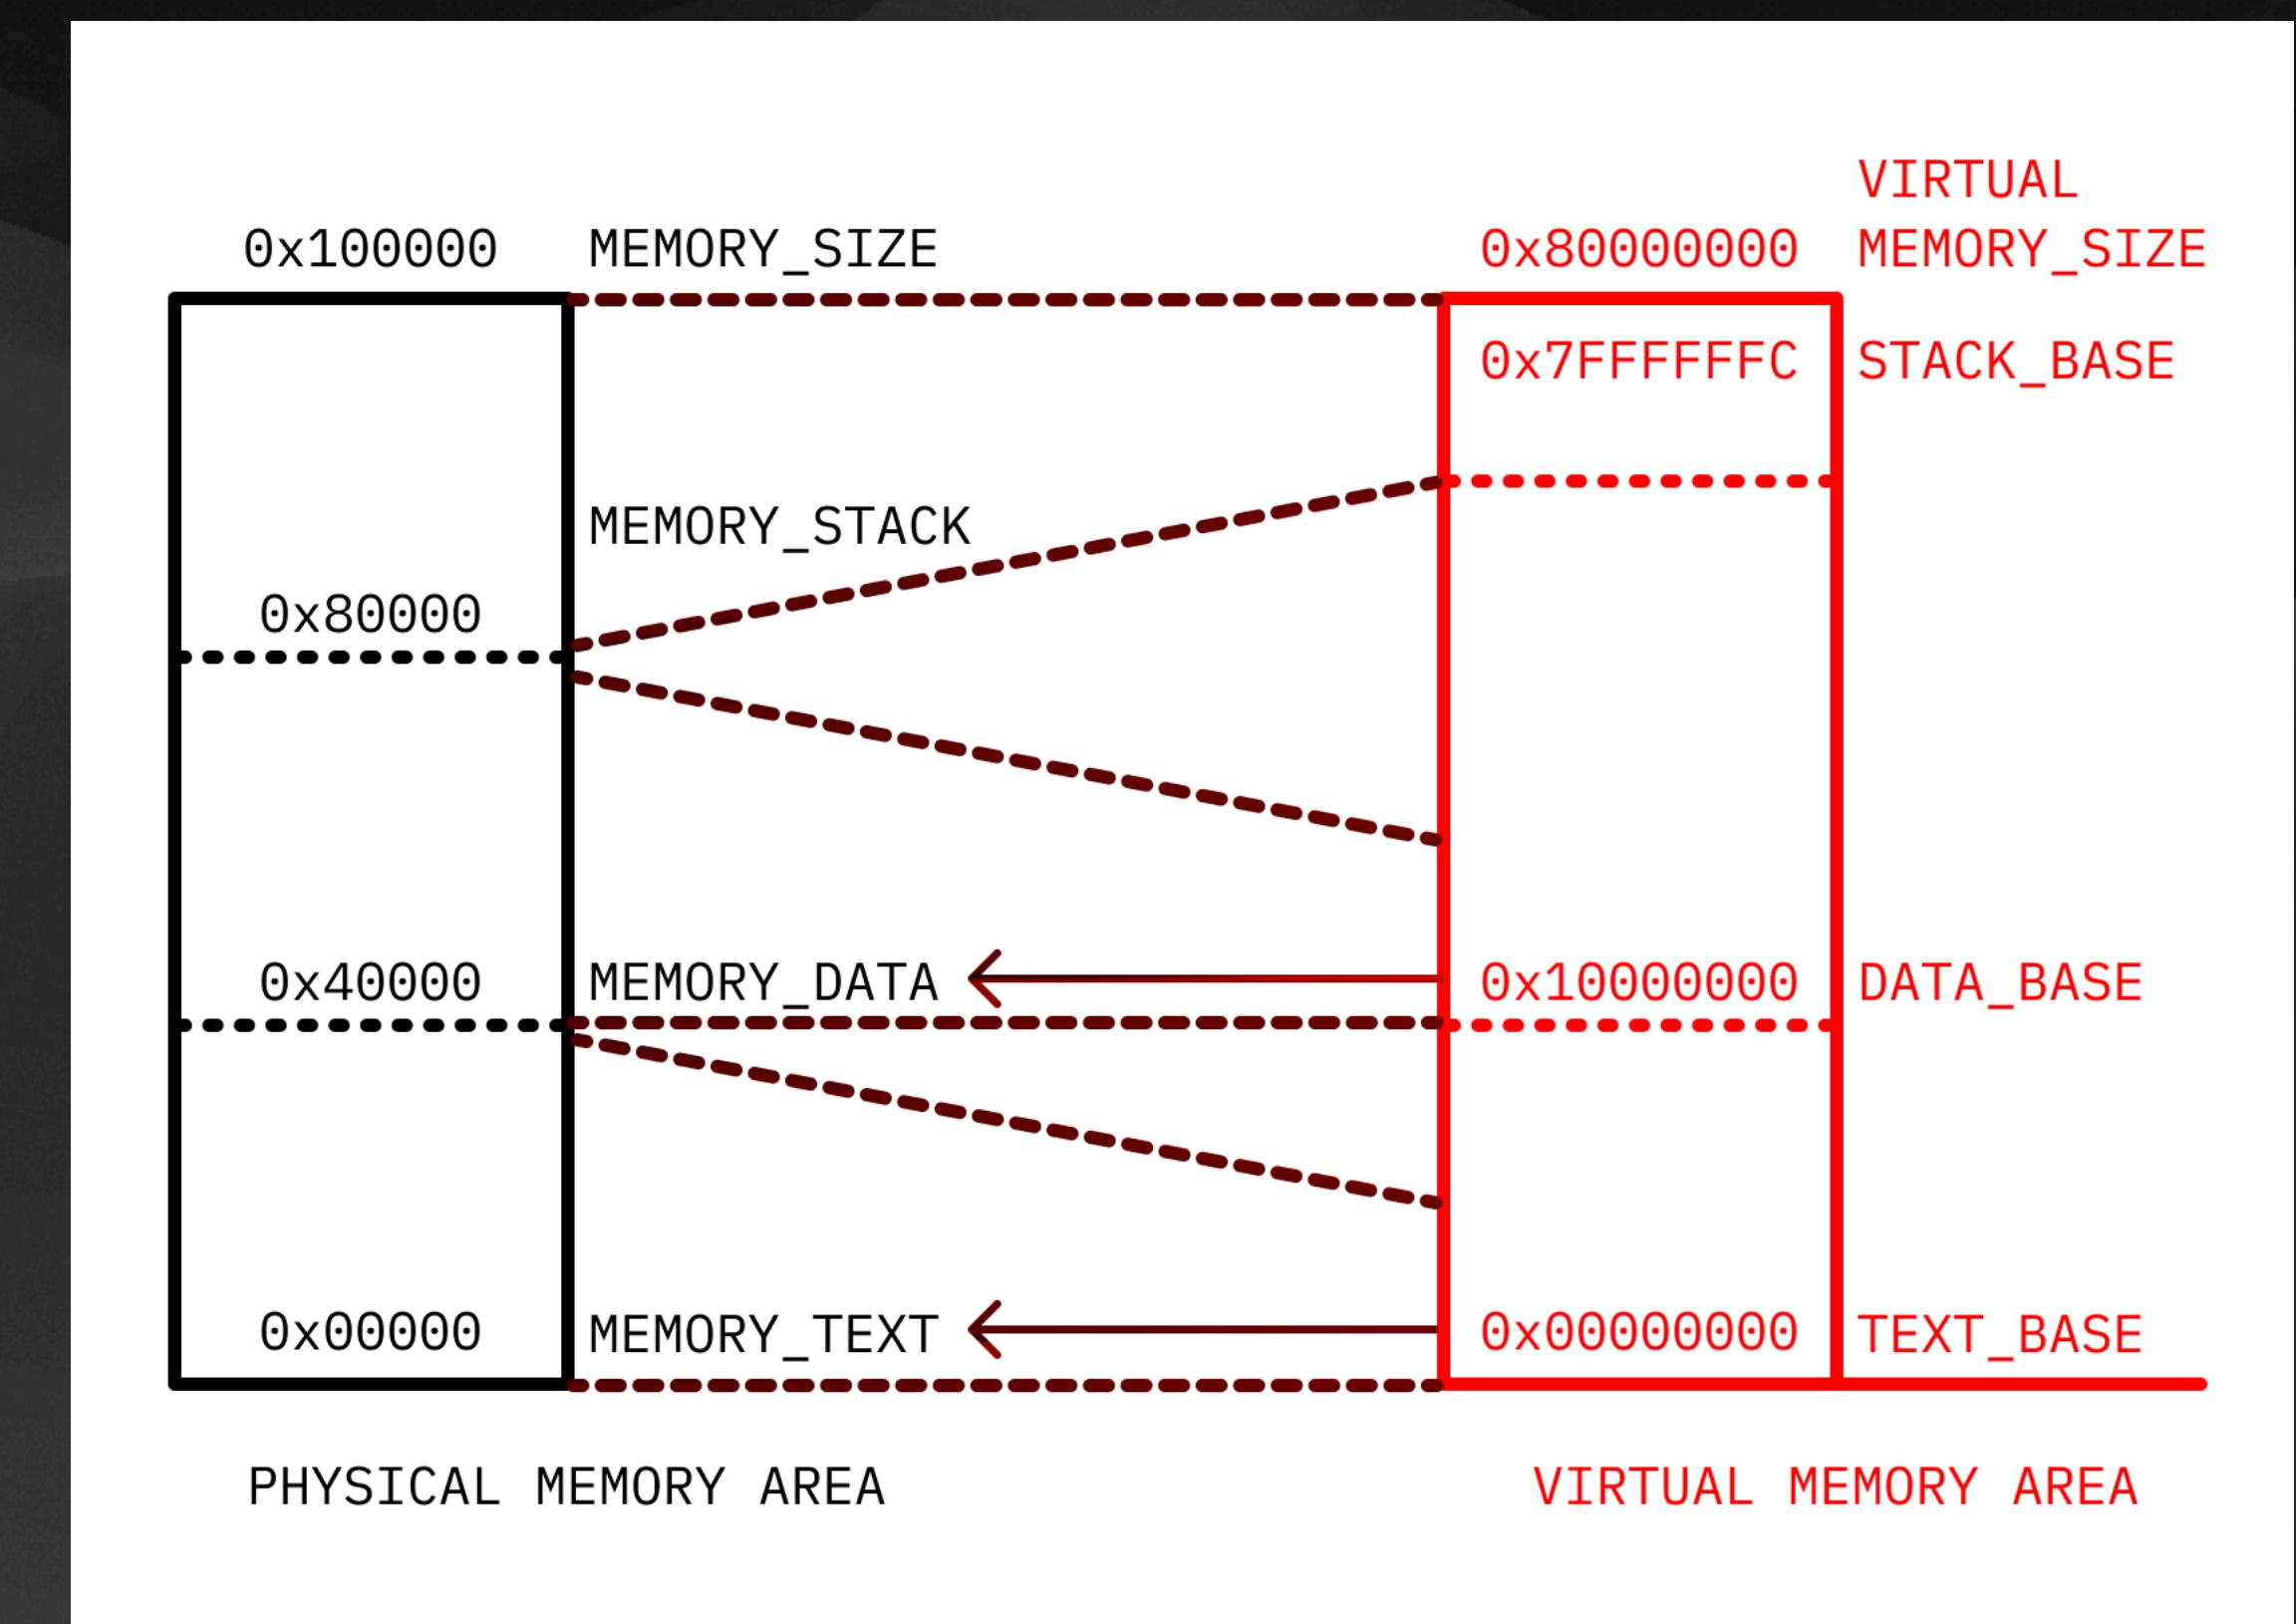
Suting Chen



# Virtual Memory

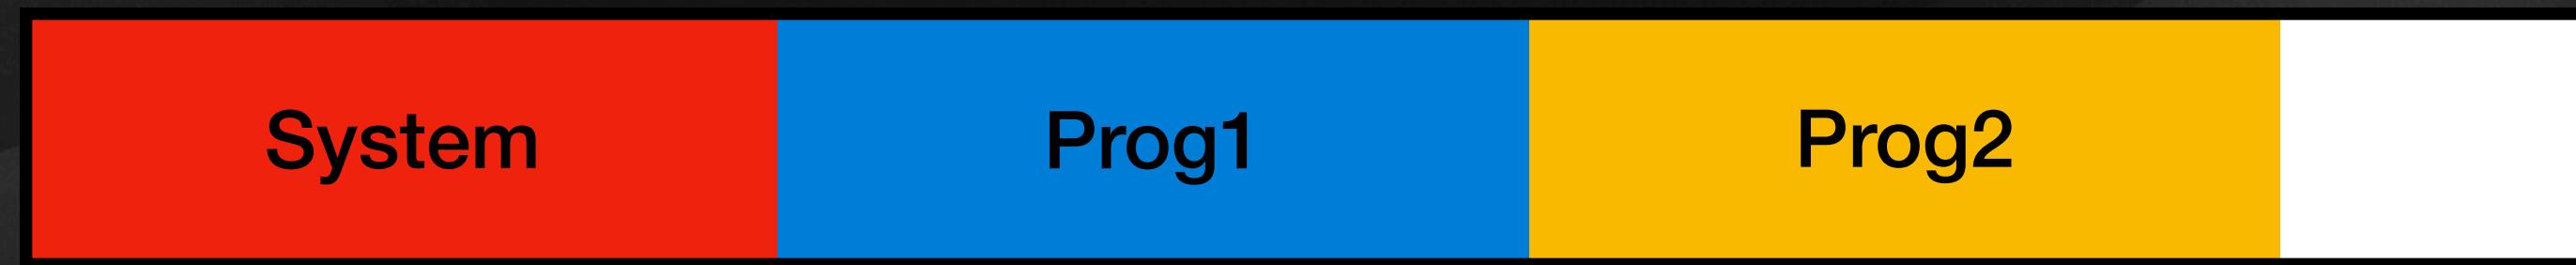
## Base and Bound

Segment	Virtual Base	Physical Base	Bound
Text	0x00000000	0x00000	0x40000
Data	0x10000000	0x40000	0x40000
stack	0x7FF80000	0x80000	0x20000
heap	-	-	-



# Virtual Memory – Base and Bound

## Memory fragmentation



$$1GB - 300MB - 300MB - 300MB = 124MB$$

Program	System	Prog1	Prog2			
Size	300M	300M	300M			

# Virtual Memory – Base and Bound

## Memory fragmentation

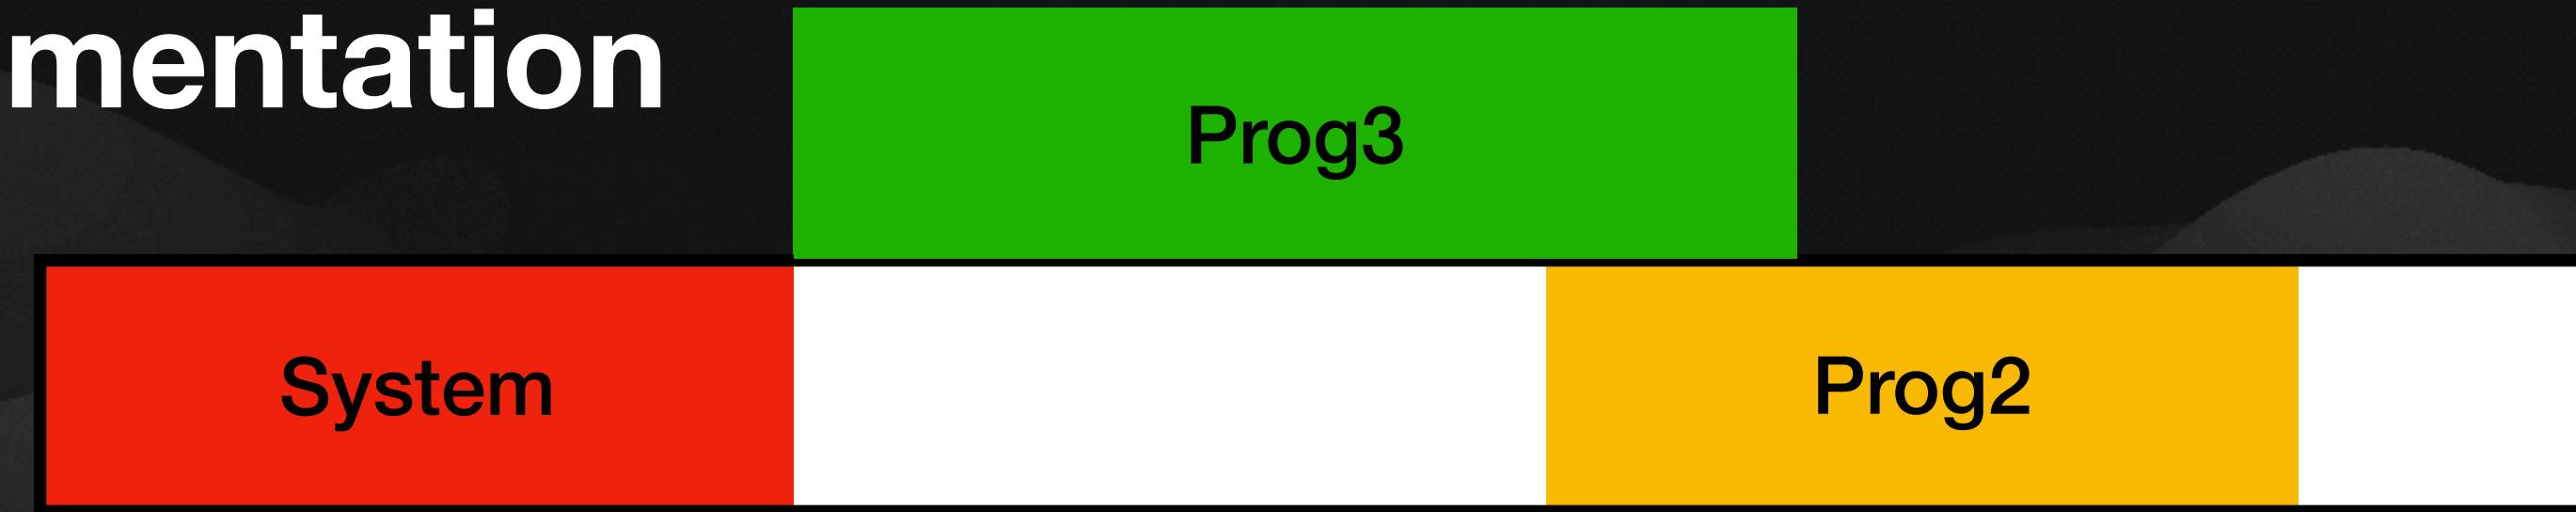


$1GB - 300MB - 300MB = 424MB$

Program	System	Prog1	Prog2			
Size	300M	-	300M			

# Virtual Memory – Base and Bound

## Memory fragmentation

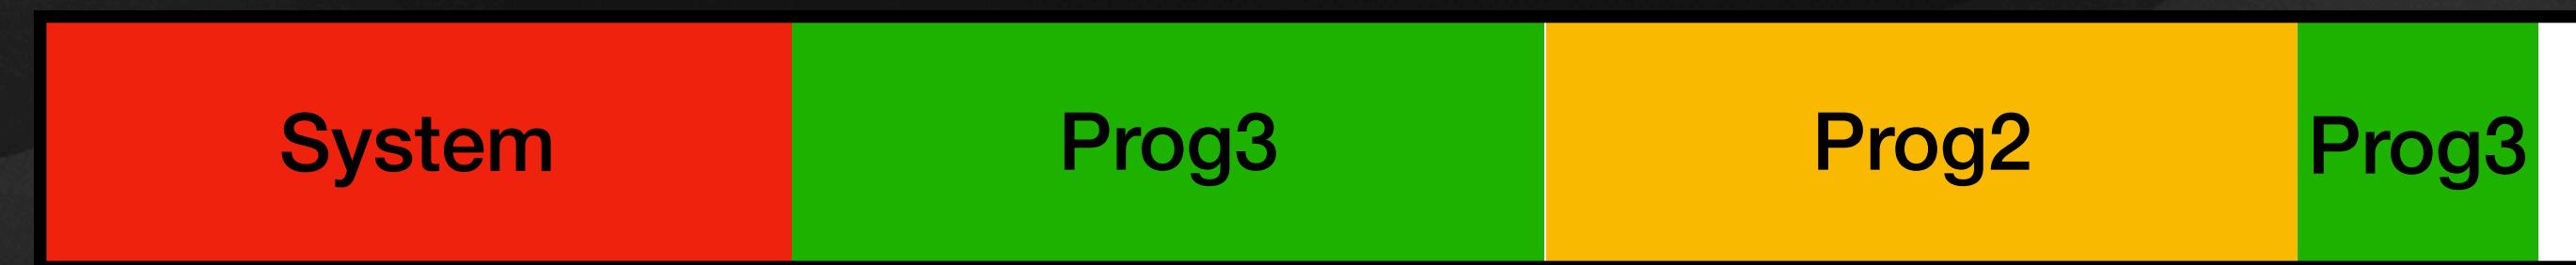


$$1GB - 300MB - 300MB - 400MB = 24MB$$

Program	System	Prog1	Prog2	Prog3		
Size	300M	-	300M	400M		

# Virtual Memory – Base and Bound

## Memory fragmentation



$$1GB - 300MB - 300MB - 400MB = 24MB$$

Program	System	Prog1	Prog2	Prog3		
Size	300M	-	300M	400M		

# Why not partition the whole device?

## Introducing Pages

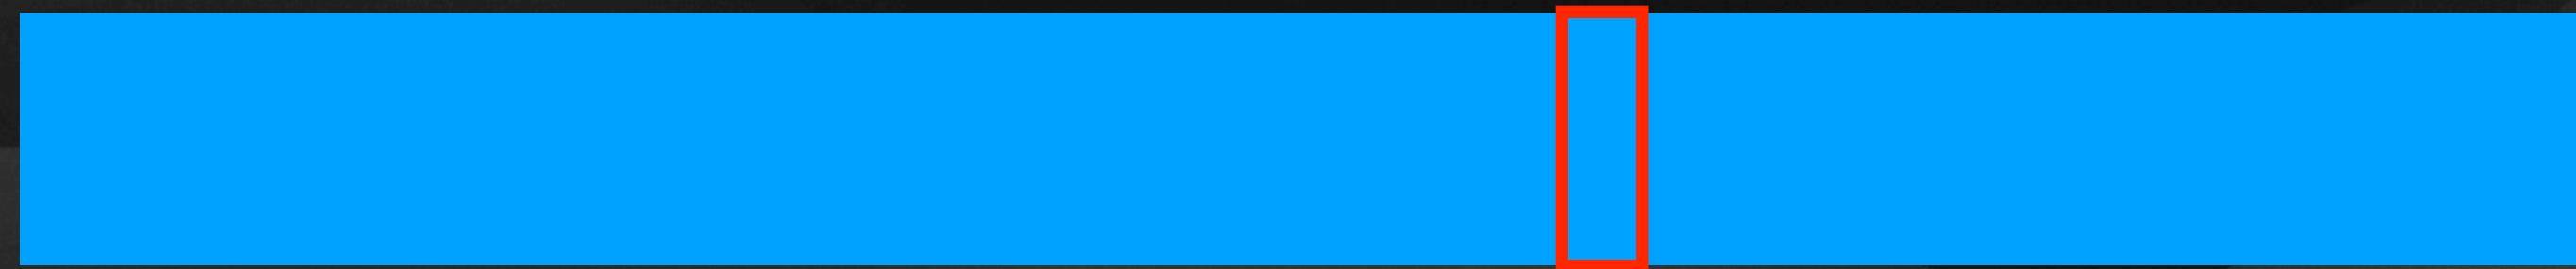
# Pages

## Something similar with cache blocks

- See memory in pages (blocks) instead of bytes
- When a process need more memory, simply allocate one page

# Pages

## On my machine



Whole memory (32GB)



getconf PAGE\_SIZE

A page (4096 Bytes)

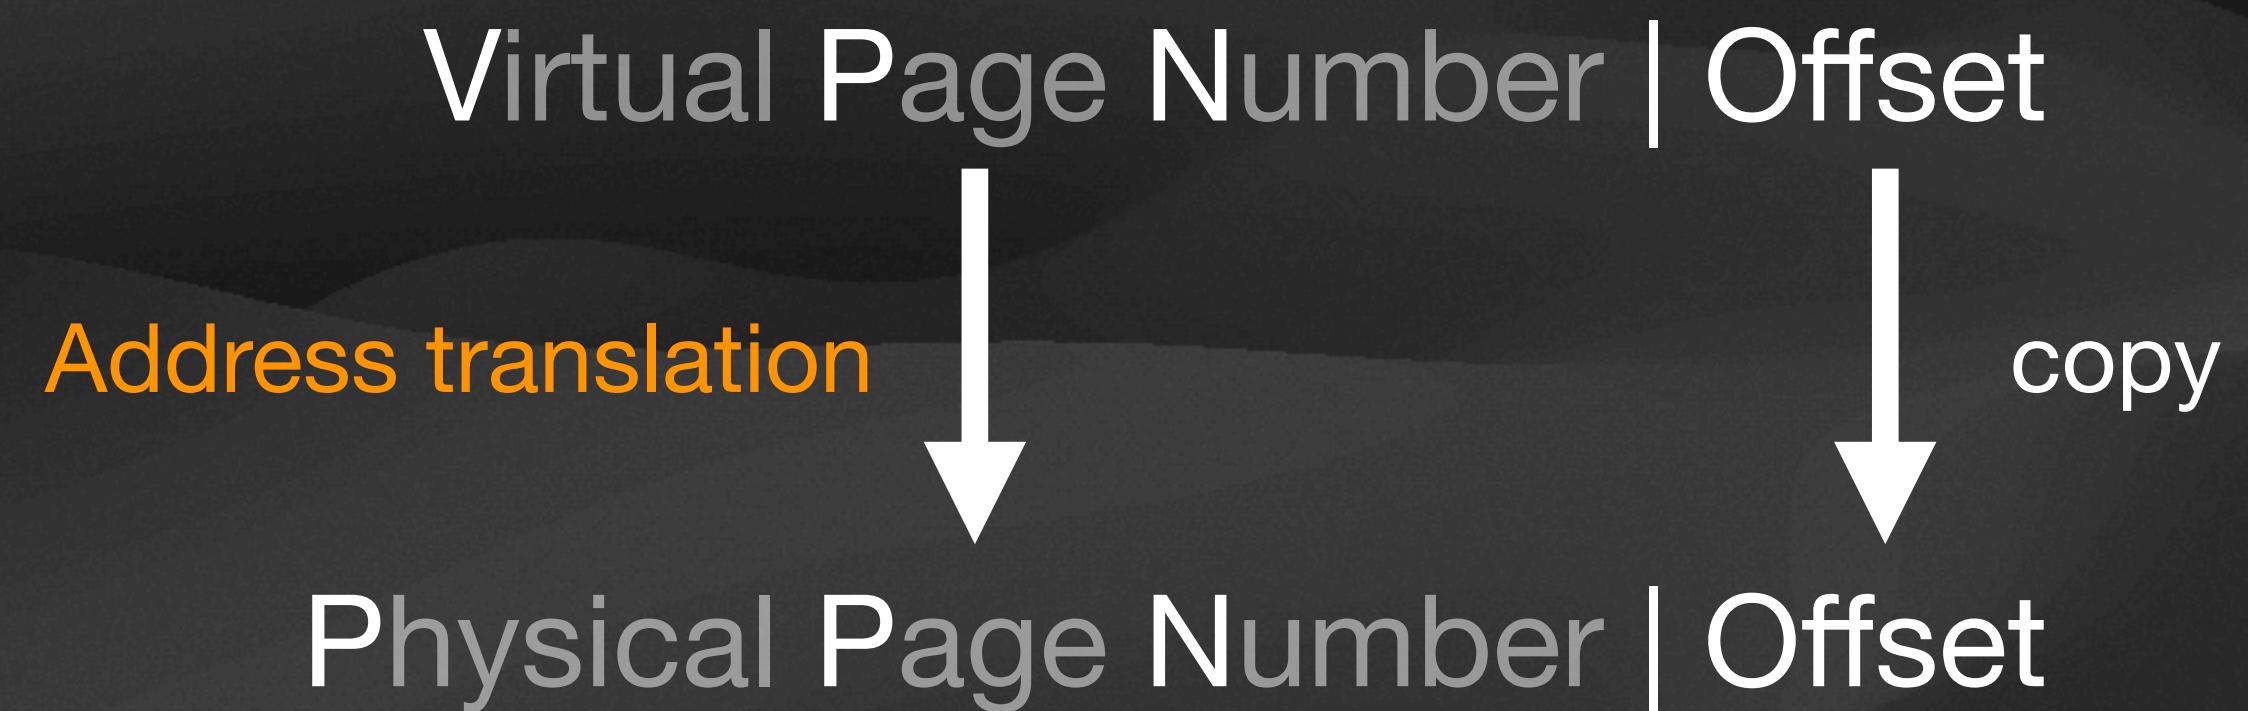


cat /proc/cpuinfo | grep cache\_alignment

A cache block (64 Bytes)

# Virtual Memory – Pages

Harness pages in **address translation**

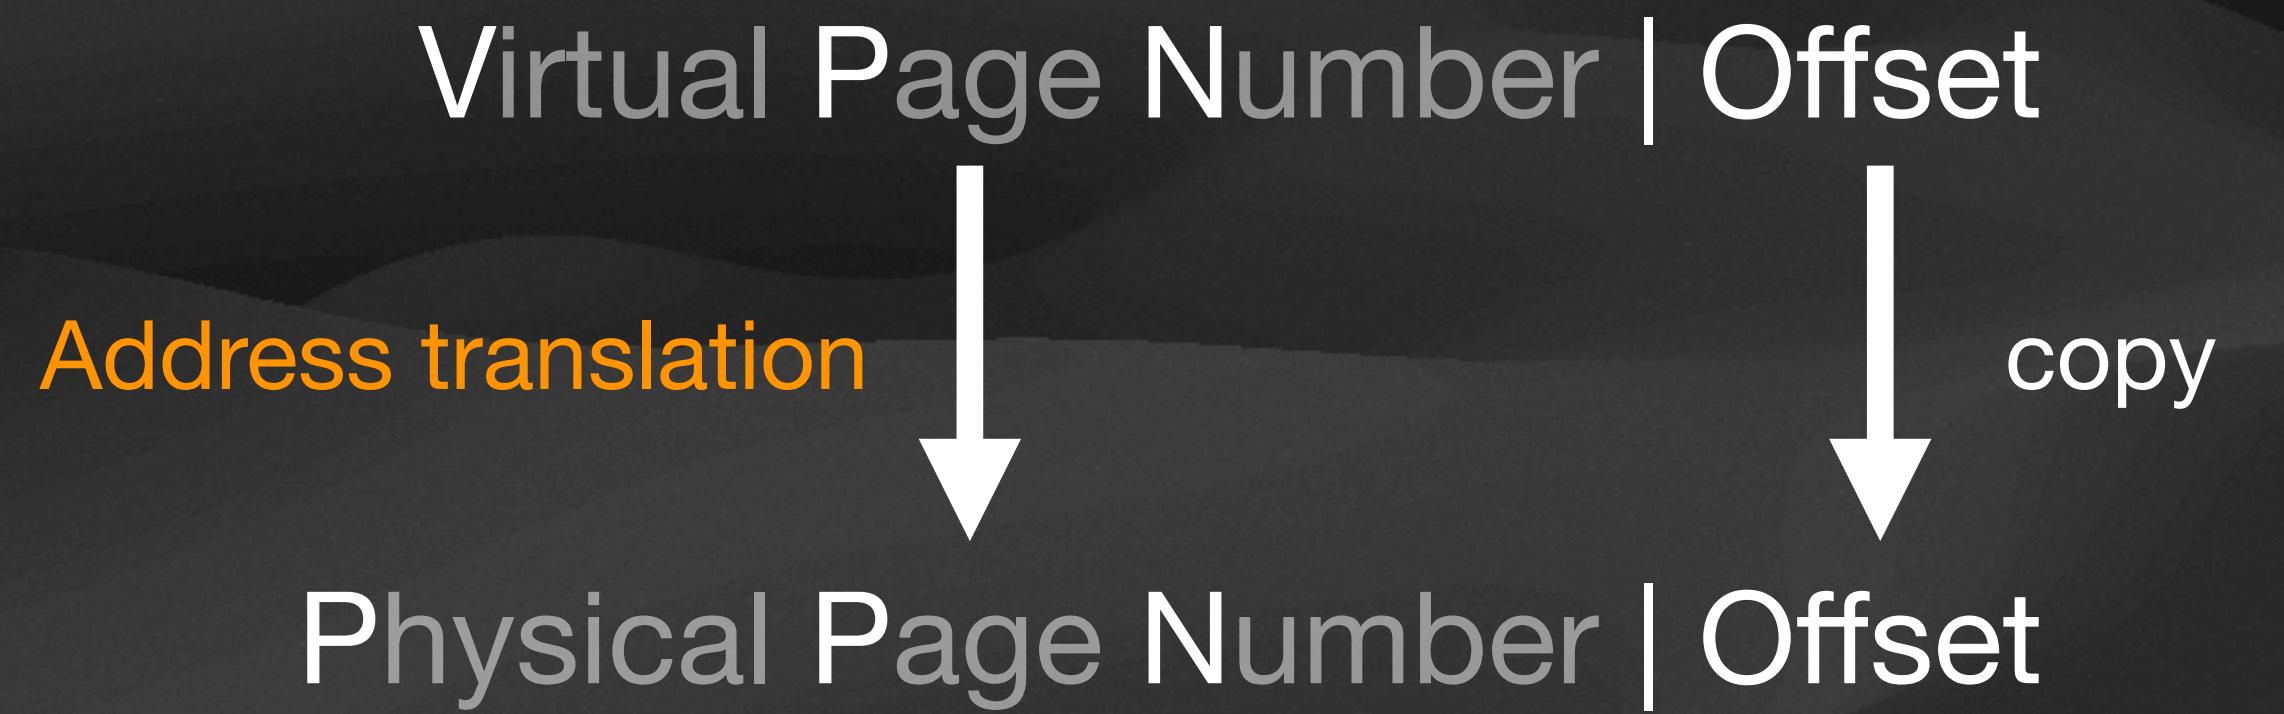


Suting Chen

Virtual	Physical
0b00000	0b0110
0b00001	0b0100
0b00010	-
0b00011	-
0b00100	0b0001
...	...
0b11111	0b0000

# Virtual Memory – Pages

Harness pages in **address translation**



Page Table Entry

**Simple Page Table**

	Physical
0	0b0110
1	0b0100
2	-
3	-
4	0b0001
...	...
31	0b0000

# Virtual Memory – Pages

## Why **NOT** Simple Page Table?

32-bit Address, 4096 Bytes page, 4 Bytes PTE

- Calculate the size of simple page table:
  1.  $\# \text{ of offset bits} = \log_2 4096 = 12$
  2.  $\# \text{ of bits of virtual page number} = 32 - 12 = 20$
  3.  $\# \text{ of PTE} = 2^{20}$
  4.  $\text{Size of Page Table} = 4 \times 2^{20} = 2^{22} \text{ Bytes} = 4MB$

# Virtual Memory – Pages

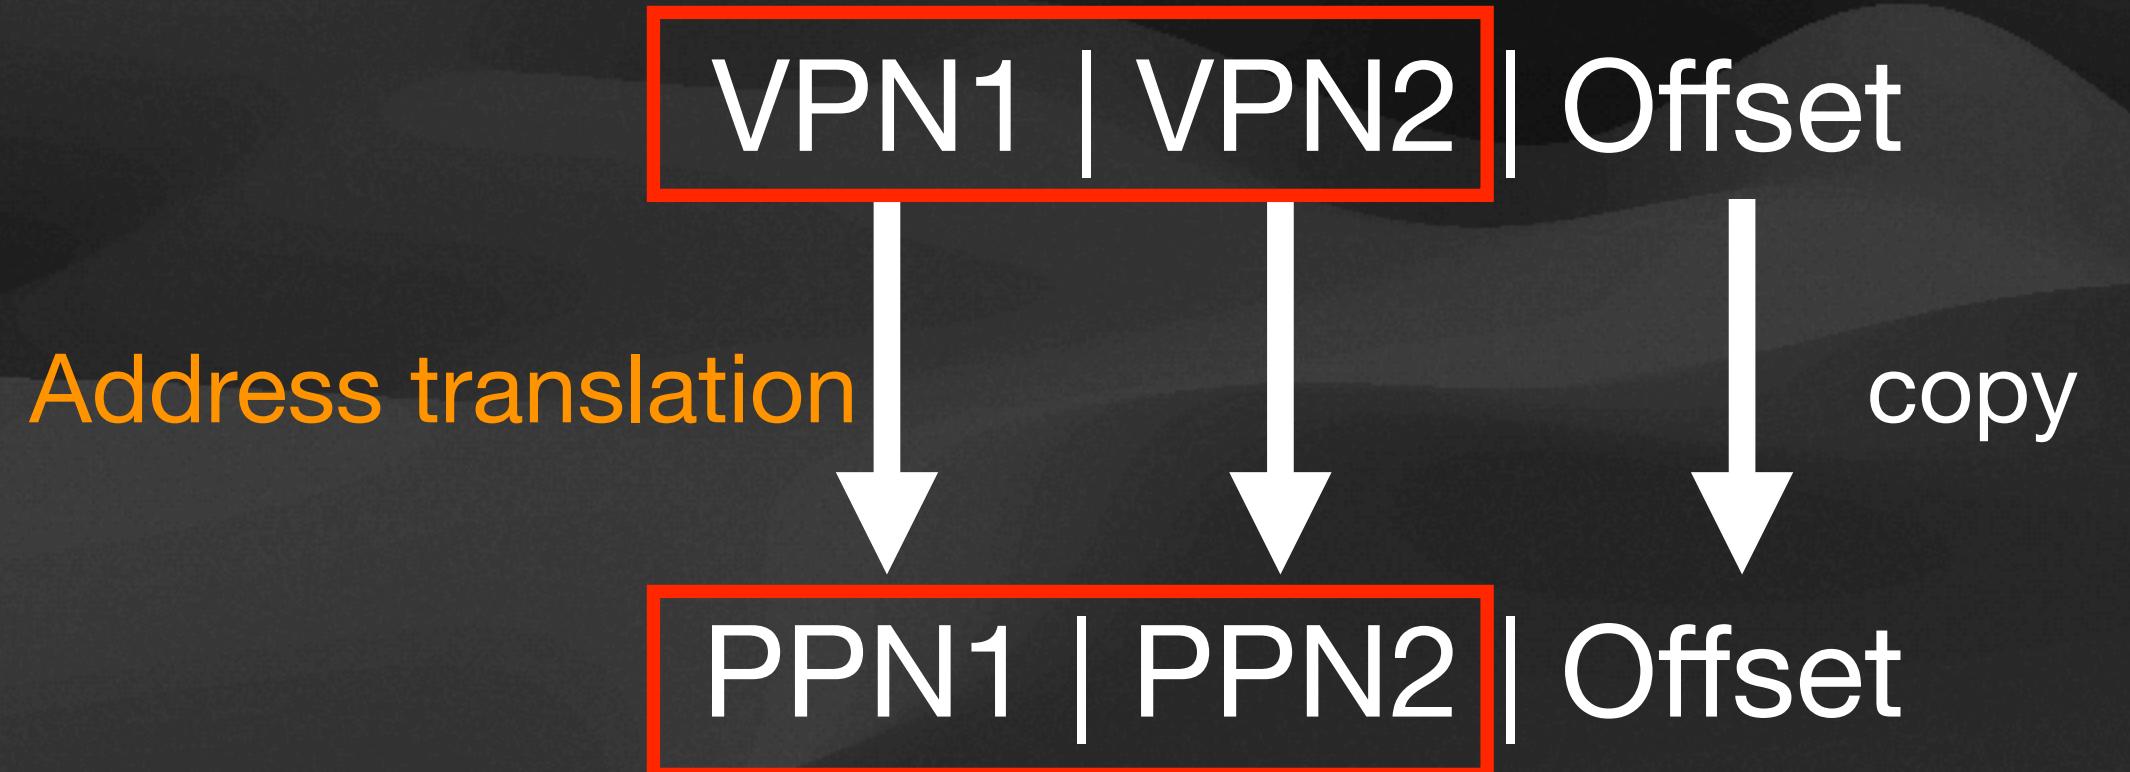
## Why **NOT** Simple Page Table?

48-bit Address, 4096 Bytes page, 4 Bytes PTE

- Calculate the size of simple page table:
  1.  $\# \text{ of offset bits} = \log_2 4096 = 12$
  2.  $\# \text{ of bits of virtual page number} = 48 - 12 = 36$
  3.  $\# \text{ of PTE} = 2^{36}$
  4.  $\text{Size of Page Table} = 4 \times 2^{36} = 2^{38} \text{ Bytes} = 256GB$

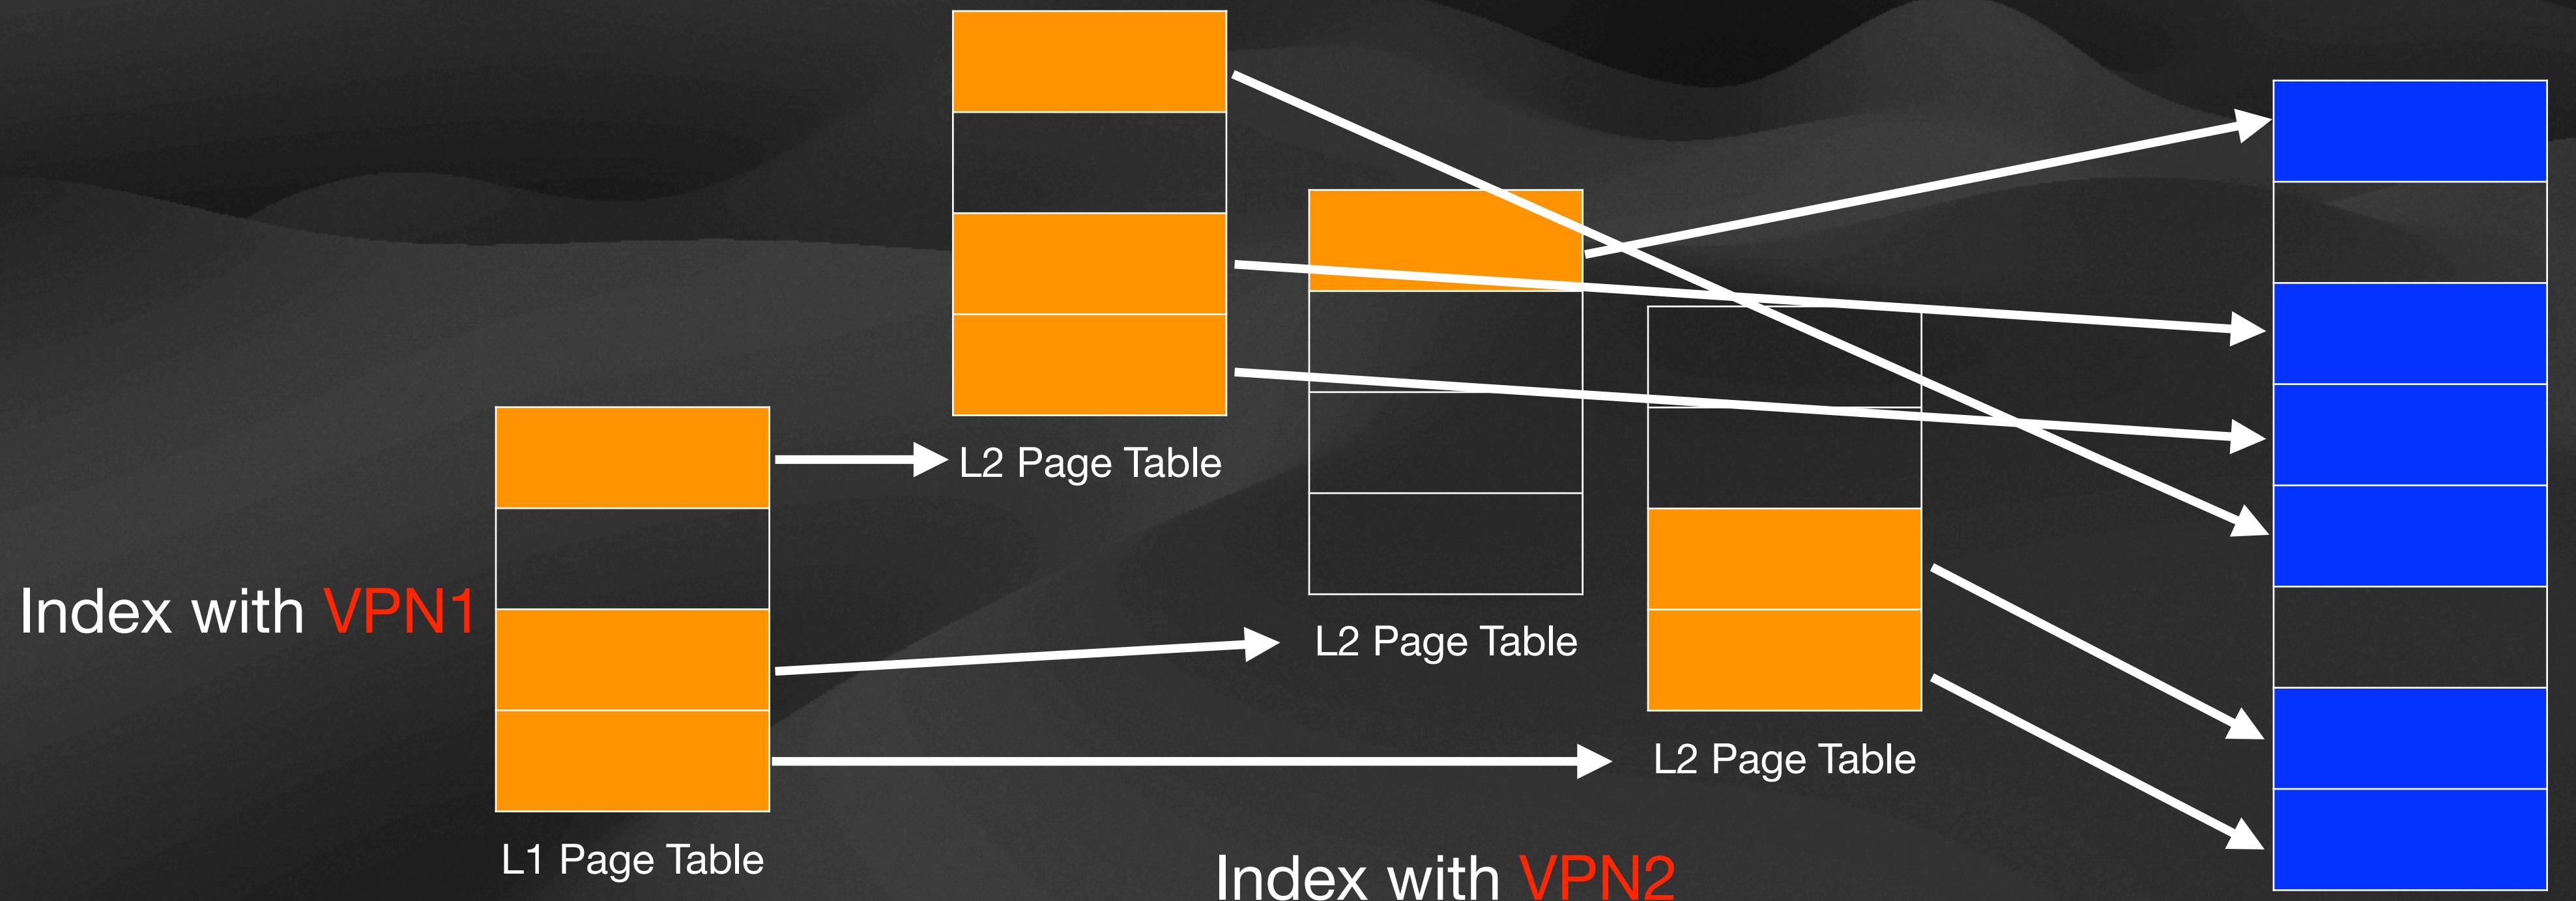
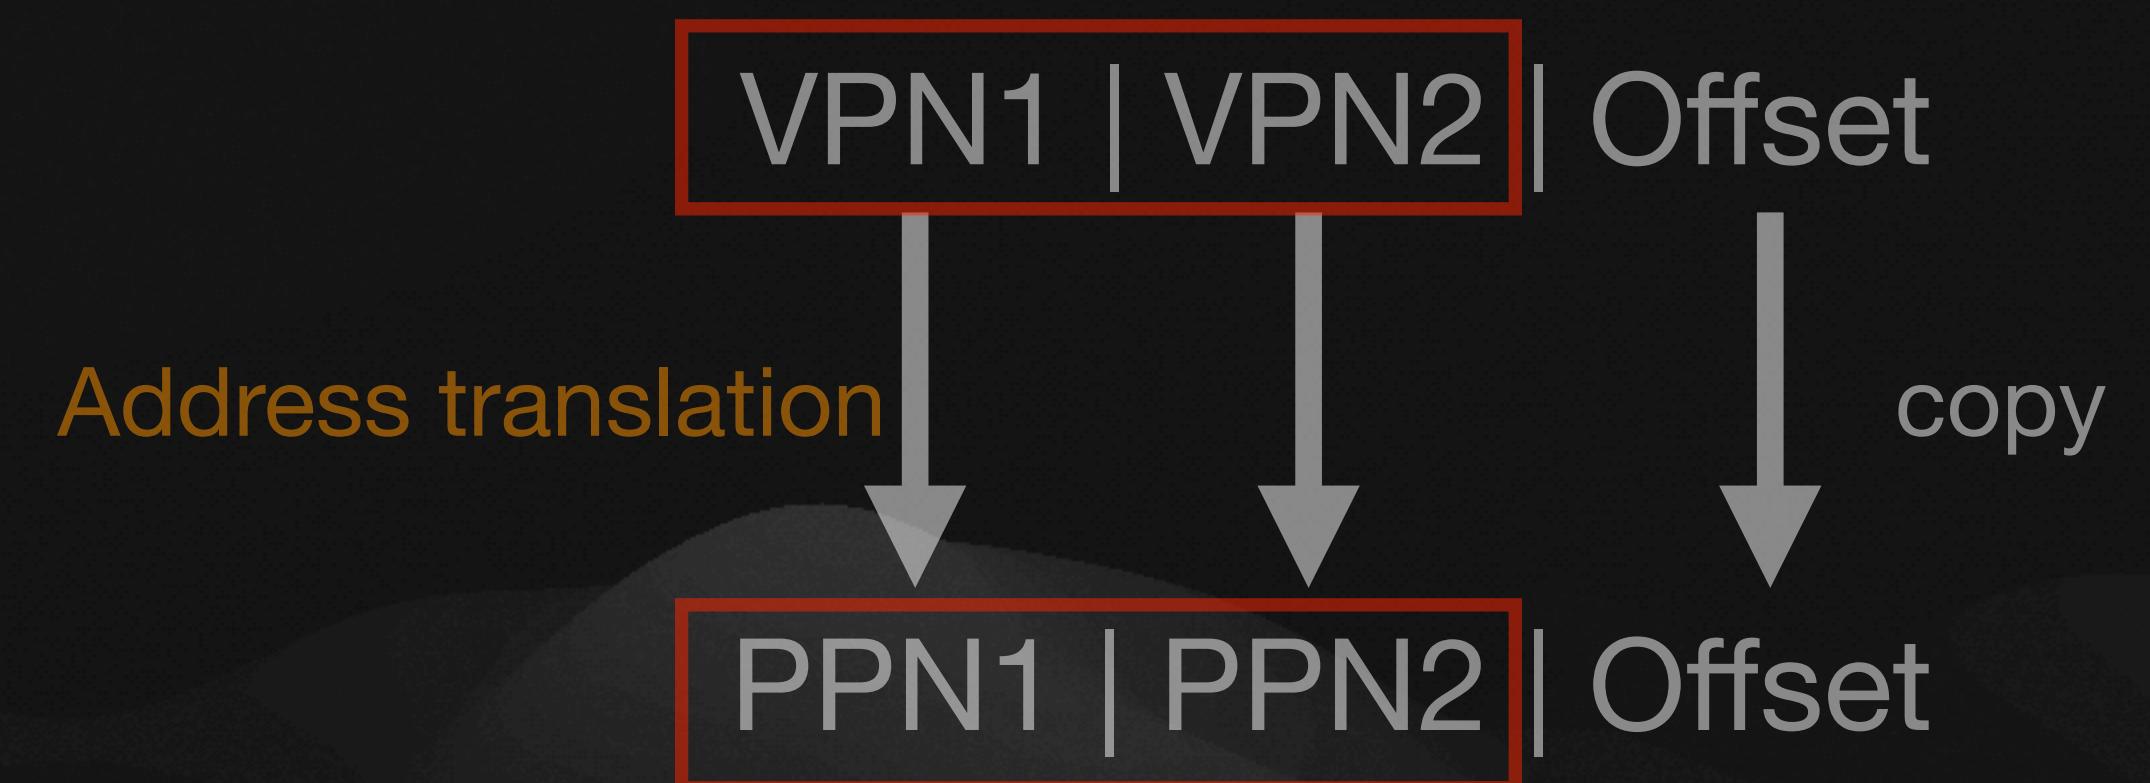
# Virtual Memory – Pages

We add hierarchy to Page Table



# Virtual Memory – Pages

We add hierarchy to Page Table



# Virtual Memory – Pages

Let's do some easy calculation!

- Given length of each segment – VPN1: 10 | VPN2: 10 | Offset: 12
  - Size of page =  $2^{12}$  Bytes = 4096 Bytes*
  - # of PTE in a L2 page table =  $2^{10}$*
  - Size of a L2 page table =  $2^{10} \times 4 = 2^{12} = 4096$  Bytes*

# Exercise

- Physical memory : 8GB
- Page size : 8KB
- Virtual address : 46-bit
- PTE : 4B
- Assume every page table exactly fits into a single page.
- How many levels of page tables would be required.

# Virtual Memory – Translation Lookaside Buffer

Recall: Memory & (fully associative) Cache

	Data
0	0x34
1	0x7c
2	0x19
3	0x01
4	0x99
...	...
31	0x12

Suting Chen

Tag	Data
01101101	0x91e989af
01000010	0xf4958104
11111011	0x82a0f265
01100101	0xfcfc4ceef
00010101	0x88b5af65
10110101	0x190349ff
10011010	0x4358f247

# Virtual Memory – Translation Lookaside Buffer

(Simple) page table & TLB

	Physical
0	0x34
1	0x7c
2	0x19
3	0x01
4	0x99
...	...
31	0x12

Suting Chen

Virtual	Physical
0b00000	0b0110
0b00001	0b0100
0b00010	-
0b00011	-
0b00100	0b0001
...	...
0b11111	0b0000

# Virtual Memory – Translation Lookaside Buffer

- Fully associative most of the time
- Can also apply replacement policy
- Each program has different address space, so invalidate TLB when switching
- TLB reach: Assume simple page table, 64 entries TLB, 4KB page
  - $TLB\ reach = 64 \times 4\ KB = 256\ KB$

# Remainders

1. P&H (our textbook) chapter 5
2. More on OS next week
3. There will be exercise on VM, so don't worry.
4. All my slides are [here](#)
5. No plagiarism! Very serious consequences!

Suting Chen